

Lab 26 - k-Nearest Neighbors classifier 2

We will continue using the Titanic training and test data from [Kaggle](https://www.kaggle.com/c/titanic) (<https://www.kaggle.com/c/titanic>) from Lab 24 and 25.

First import the necessary libraries.

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
%matplotlib inline
```

Loading and cleaning the data

The code for loading in the data and cleaning it from Lab 25 is below.

```
In [2]: train = pd.read_csv("train.csv")
train.head()
```

Out[2]:

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	C85	C
3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

```
In [3]: # fill in missing age data
train["Age"] = train["Age"].fillna(train["Age"].median())

# fill in the missing embarked data
train["Embarked"] = train["Embarked"].fillna("S")
```

```
In [4]: # create dummy variables for passenger class, sex, and embarked
train2 = pd.get_dummies(train, columns = ["Pclass", "Sex", "Embarked"], drop_first = True)
train2.head()
```

Out[4]:

	PassengerId	Survived	Name	Age	SibSp	Parch	Ticket	Fare	Cabin	Pclass_2	P
0	1	0	Braund, Mr. Owen Harris	22.0	1	0	A/5 21171	7.2500	NaN	0	
1	2	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	38.0	1	0	PC 17599	71.2833	C85	0	
2	3	1	Heikkinen, Miss. Laina	26.0	0	0	STON/O2. 3101282	7.9250	NaN	0	
3	4	1	Futrelle, Mrs. Jacques Heath	35.0	1	0	113803	53.1000	C123	0	

```
In [5]: # remove the remaining qualitative columns
train2.drop("Cabin",axis = 1,inplace = True)
train2.drop("Name",axis = 1,inplace = True)
train2.drop("Ticket",axis = 1,inplace = True)

# we should also drop PassengerId, although we did not last lab
train2.drop("PassengerId",axis = 1, inplace = True)
```

```
In [6]: # split the original training data into training and test sets
X_train,X_test,y_train, y_test =train_test_split(train2.drop("Survived",axis=1),train2["Survived"],
                                                test_size = 0.2)
```

```
In [7]: # create a 3-nearest neighbor classifier
knn = KNeighborsClassifier(n_neighbors=3)
# fit the classifier to the training data
knn.fit(X_train, y_train)
# test and score the classifier on our test data (part of the original training data)
# notice this line corrects a mistake in Lab 25
knn.score(X_test, y_test)
```

Out[7]: 0.7262569832402235

Now we are going to try running our classifier on the Kaggle test data and use all of our training data to fit the classifier.

First, load the test data from Kaggle into the dataframe `test` .

```
In [8]: test = pd.read_csv("test.csv")
test.head()
```

Out[8]:

	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	892	3	Kelly, Mr. James	male	34.5	0	0	330911	7.8292	NaN	Q
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female	47.0	1	0	363272	7.0000	NaN	S
2	894	2	Myles, Mr. Thomas Francis	male	62.0	0	0	240276	9.6875	NaN	Q
3	895	3	Wirz, Mr. Albert	male	27.0	0	0	315154	8.6625	NaN	S
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	22.0	1	1	3101298	12.2875	NaN	S

We have to process the test data in the same way as the training data, namely filling in the missing age and embarked data, creating the dummy variables for Pclass, Sex, and Embarked, and dropping the Cabin, Name, and Ticket columns. Do this below, adding as many extra cells as you need.

```
In [9]: test["Age"] = test["Age"].fillna(test["Age"].median())
test["Embarked"] = test["Embarked"].fillna("S")

test2 = pd.get_dummies(test, columns = ["Pclass", "Sex", "Embarked"], drop_first = True)

test2.drop("Cabin", axis = 1, inplace = True)
test2.drop("Name", axis = 1, inplace = True)
test2.drop("Ticket", axis = 1, inplace = True)

test2.head()
```

Out[9]:

	PassengerId	Age	SibSp	Parch	Fare	Pclass_2	Pclass_3	Sex_male	Embarked_Q	Embarked
0	892	34.5	0	0	7.8292	0	1	1	1	
1	893	47.0	1	0	7.0000	0	1	0	0	
2	894	62.0	0	0	9.6875	1	0	1	1	
3	895	27.0	0	0	8.6625	0	1	1	0	
4	896	22.0	1	1	12.2875	0	1	0	0	

Store the column `PassengerId` in a variable. We need this information for submitting our predictions

```
In [10]: passengerId = test2["PassengerId"]
passengerId
```

```
Out[10]: 0      892
1      893
2      894
3      895
4      896
...
413    1305
414    1306
415    1307
416    1308
417    1309
Name: PassengerId, Length: 418, dtype: int64
```

Drop the `PassengerId` column from the test data.

```
In [11]: test2.drop("PassengerId",axis = 1, inplace = True)
```

```
In [12]: test2.head()
```

```
Out[12]:
```

	Age	SibSp	Parch	Fare	Pclass_2	Pclass_3	Sex_male	Embarked_Q	Embarked_S
0	34.5	0	0	7.8292	0	1	1	1	0
1	47.0	1	0	7.0000	0	1	0	0	1
2	62.0	0	0	9.6875	1	0	1	1	0
3	27.0	0	0	8.6625	0	1	1	0	1
4	22.0	1	1	12.2875	0	1	0	0	1

Next we split up our training data into the answer (the `Survived` column) and the input data (all other columns).

First, store the `Survived` column in the variable `y_train_kaggle` .

```
In [13]: y_train_kaggle = train2["Survived"]
y_train_kaggle.head()
```

```
Out[13]: 0    0
1    1
2    1
3    1
4    0
Name: Survived, dtype: int64
```

Next, drop the `Survived` column from the training data, and store the new data frame in the variable `X_train_kaggle` .

```
In [14]: X_train_kaggle = train2.drop("Survived",axis = 1)
X_train_kaggle.head()
```

Out[14]:

	Age	SibSp	Parch	Fare	Pclass_2	Pclass_3	Sex_male	Embarked_Q	Embarked_S
0	22.0	1	0	7.2500	0	1	1	0	1
1	38.0	1	0	71.2833	0	0	0	0	0
2	26.0	0	0	7.9250	0	1	0	0	1
3	35.0	1	0	53.1000	0	0	0	0	1
4	35.0	0	0	8.0500	0	1	1	0	1

Create a new k-nearest neighbors object and fit it on the entire training data (X_train_kaggle).

```
In [15]: knn_kaggle = KNeighborsClassifier(n_neighbors=3)
knn_kaggle.fit(X_train_kaggle,y_train_kaggle)
```

Out[15]: KNeighborsClassifier(n_neighbors=3)

Now use our fitted classifier to make predictions on the test data (test2), and store it in the variable y_pred .

```
In [17]: y_pred = knn_kaggle.predict(test2)
```

```
-----  
ValueError                                Traceback (most recent call last)  
C:\Users\S1406~1\ALI\AppData\Local\Temp\ipykernel_18680\290278459.py in <mod  
ule>  
----> 1 y_pred = knn_kaggle.predict(test2)  
  
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\neighbors\_classification.py in  
predict(self, X)  
    193         Class labels for each data sample.  
    194         """  
--> 195         X = check_array(X, accept_sparse='csr')  
    196  
    197         neigh_dist, neigh_ind = self.kneighbors(X)  
  
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in inner_f(*a  
rgs, **kwargs)  
    61         extra_args = len(args) - len(all_args)  
    62         if extra_args <= 0:  
--> 63             return f(*args, **kwargs)  
    64  
    65         # extra_args > 0  
  
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_arra  
y(array, accept_sparse, accept_large_sparse, dtype, order, copy, force_all_finite, ensur  
e_2d, allow_nd, ensure_min_samples, ensure_min_features, estimator)  
    718  
    719         if force_all_finite:  
--> 720             _assert_all_finite(array,  
    721                             allow_nan=force_all_finite == 'allow-nan')  
    722  
  
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in _assert_all  
_finite(X, allow_nan, msg_dtype)  
    101         not allow_nan and not np.isfinite(X).all()):  
    102         type_err = 'infinity' if allow_nan else 'NaN, infinity'  
--> 103         raise ValueError(  
    104             msg_err.format  
    105             (type_err,
```

ValueError: Input contains NaN, infinity or a value too large for dtype('float64').

The error message says something about NaN. Could there be missing data (NaN) in a different column in the test data? Use describe to see if this is the case.

```
In [18]: test2.describe()
```

Out[18]:

	Age	SibSp	Parch	Fare	Pclass_2	Pclass_3	Sex_male	Embarke
count	418.000000	418.000000	418.000000	417.000000	418.000000	418.000000	418.000000	418.000000
mean	29.599282	0.447368	0.392344	35.627188	0.222488	0.521531	0.636364	0.110380
std	12.703770	0.896760	0.981429	55.907576	0.416416	0.500135	0.481622	0.310000
min	0.170000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	23.000000	0.000000	0.000000	7.895800	0.000000	0.000000	0.000000	0.000000
50%	27.000000	0.000000	0.000000	14.454200	0.000000	1.000000	1.000000	0.000000
75%	35.750000	1.000000	0.000000	31.500000	0.000000	1.000000	1.000000	0.000000
max	76.000000	8.000000	9.000000	512.329200	1.000000	1.000000	1.000000	1.000000

The fare column is missing one value. Fill it in with the median fare.

```
In [19]: test2["Fare"] = test2["Fare"].fillna(test2["Fare"].median())
```

Now try making the prediction again.

```
In [20]: y_pred = knn_kaggle.predict(test2)
```

```
In [21]: y_pred
```

Out[21]: array([0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0,
1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0,
1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0,
0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1,
1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0,
1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1,
0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1,
0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0,
0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1,
0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0,
1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1,
0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0,
1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1,
1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0,
0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0,
0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1,
1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,
0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1],
dtype=int64)

Finally, we want to write our prediction to a file, along with the passenger ID.

```
In [22]: df = pd.DataFrame(data = {"PassengerId":passengerId, "Survived":y_pred})
```

```
In [23]: df.to_csv("test1.csv",index = 0)
```

```
In [ ]:
```